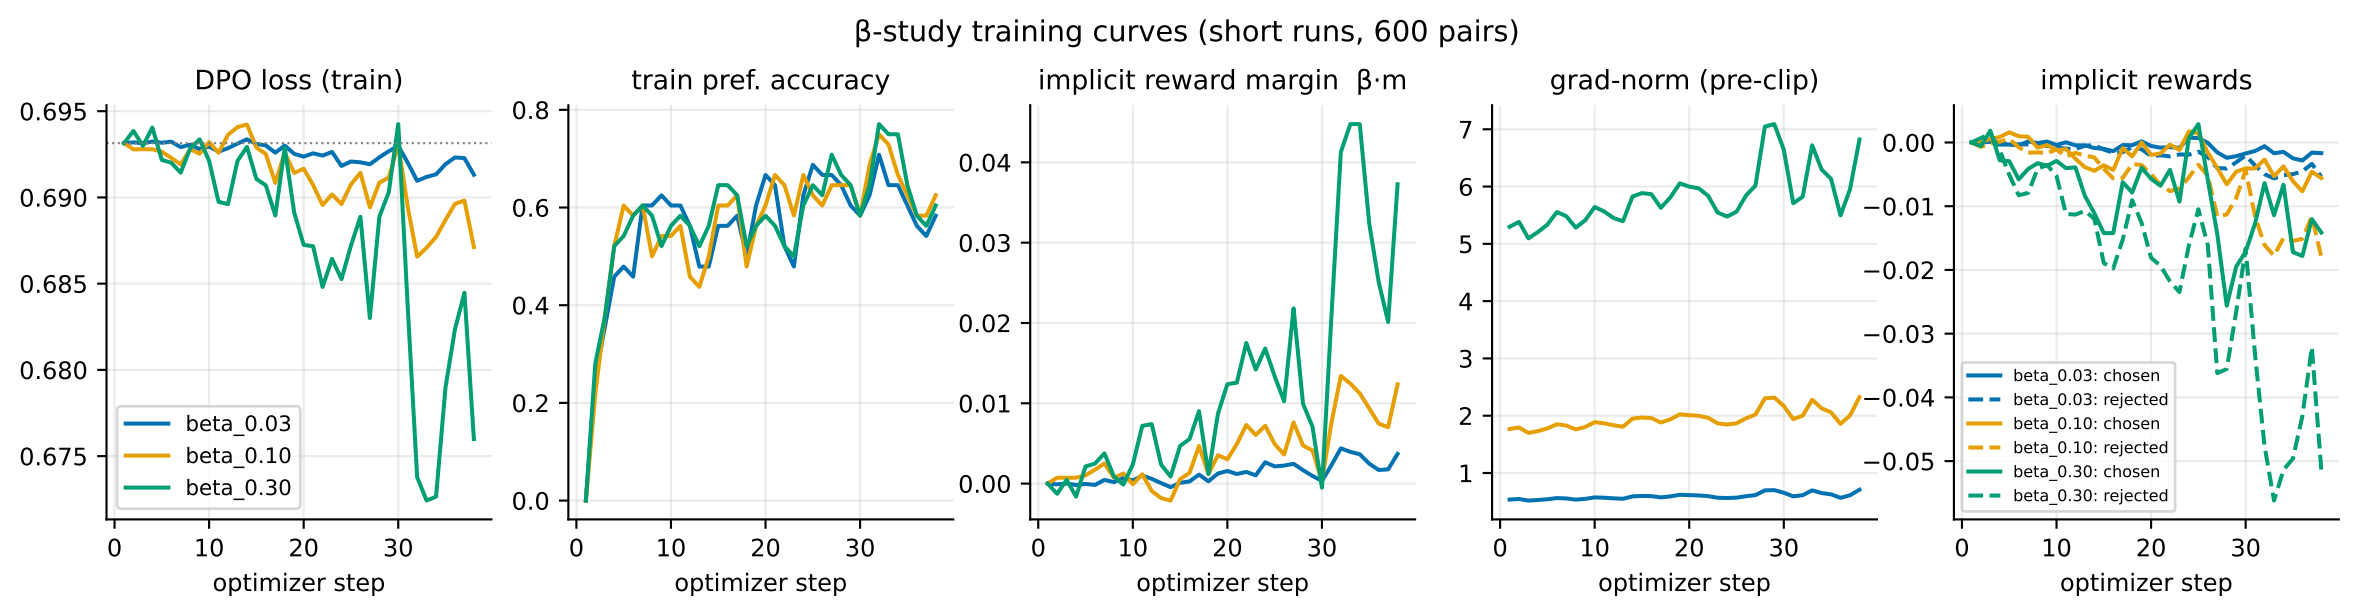
β-study training curves (short runs, 600 pairs)
DPO loss (train)
train pref. accuracy
implicit reward margin β·m
grad-norm (pre-clip)
implicit rewards
optimizer step
beta_0.03
beta_0.10
beta_0.30
beta_0.03: chosen
beta_0.03: rejected
beta_0.10: chosen
beta_0.10: rejected
beta_0.30: chosen
beta_0.30: rejected