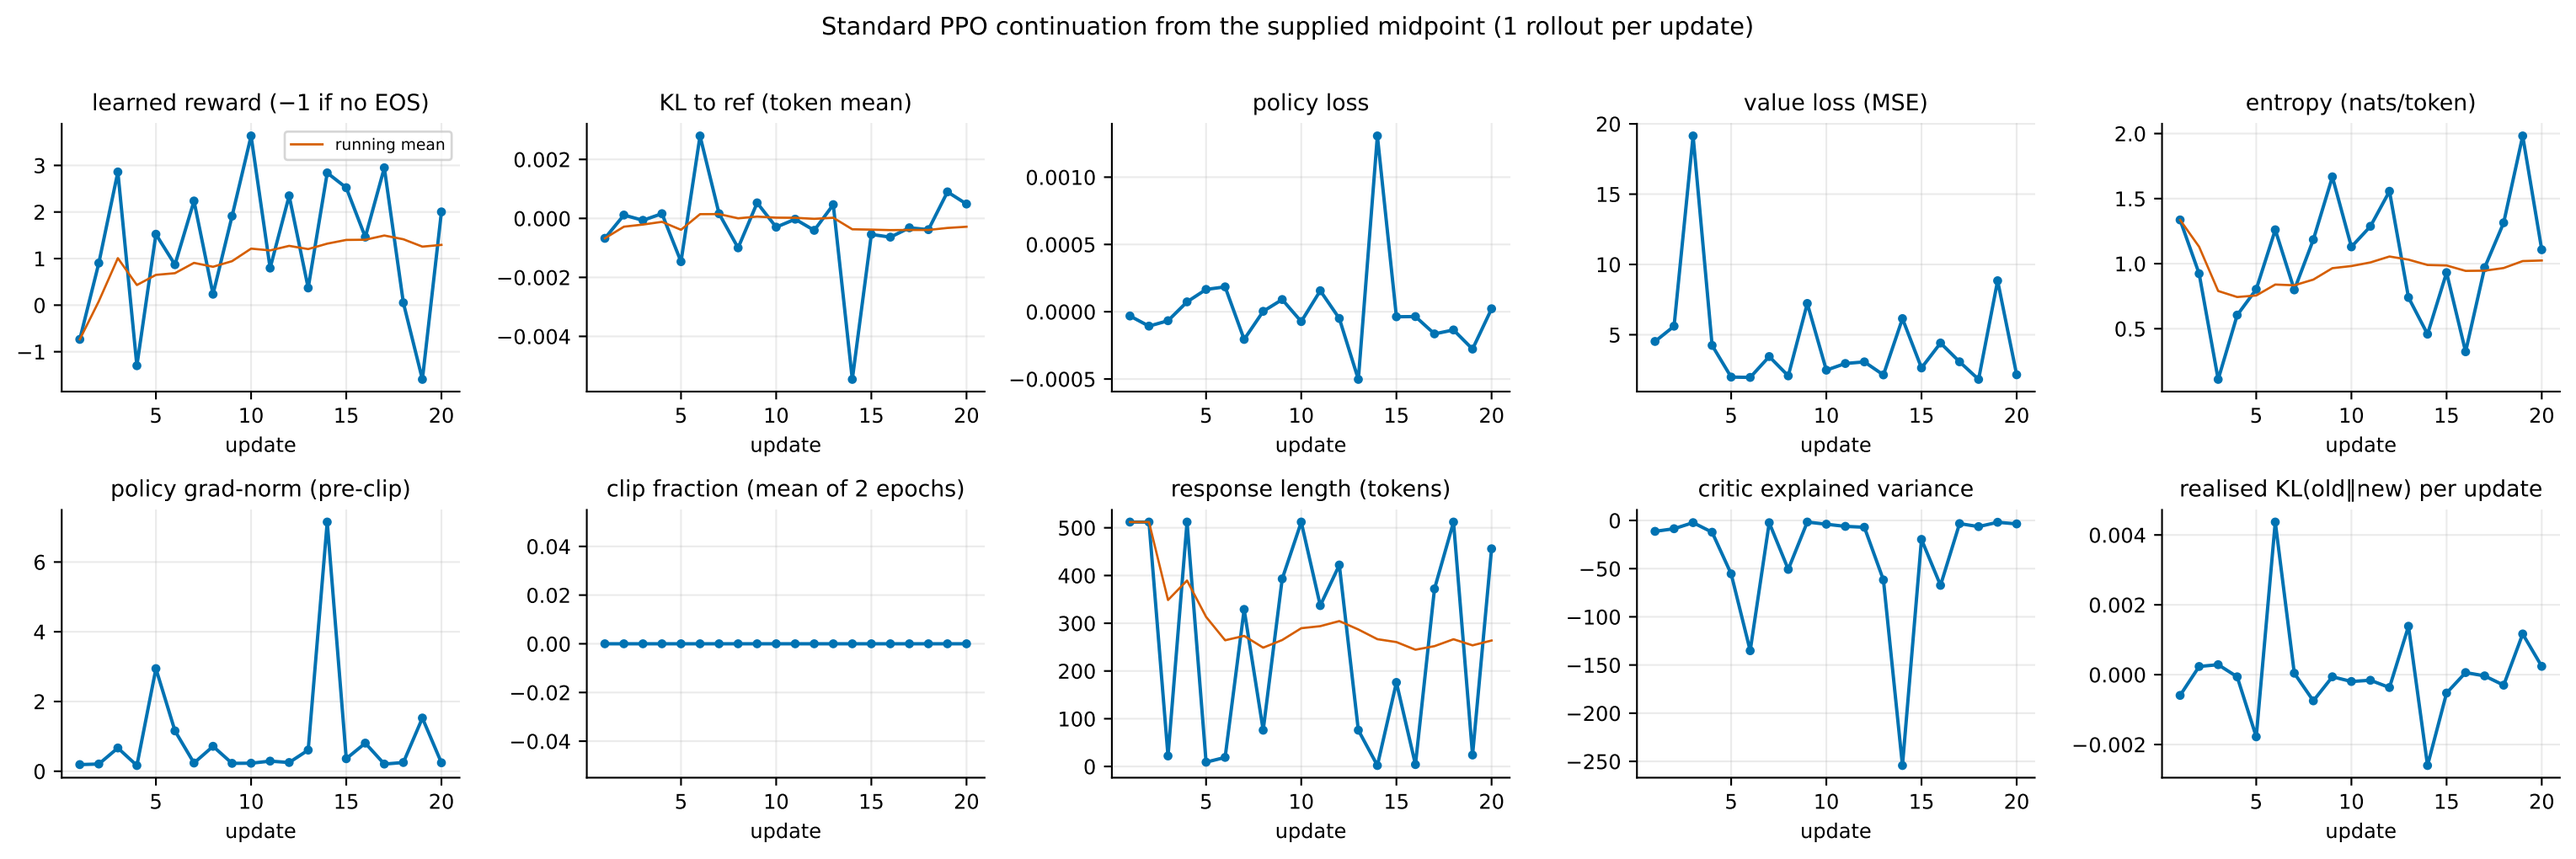
Standard PPO continuation from the supplied midpoint (1 rollout per update)
learned reward (−1 if no EOS)
running mean
KL to ref (token mean)
policy loss
value loss (MSE)
entropy (nats/token)
policy grad-norm (pre-clip)
clip fraction (mean of 2 epochs)
response length (tokens)
critic explained variance
realised KL(old‖new) per update
update