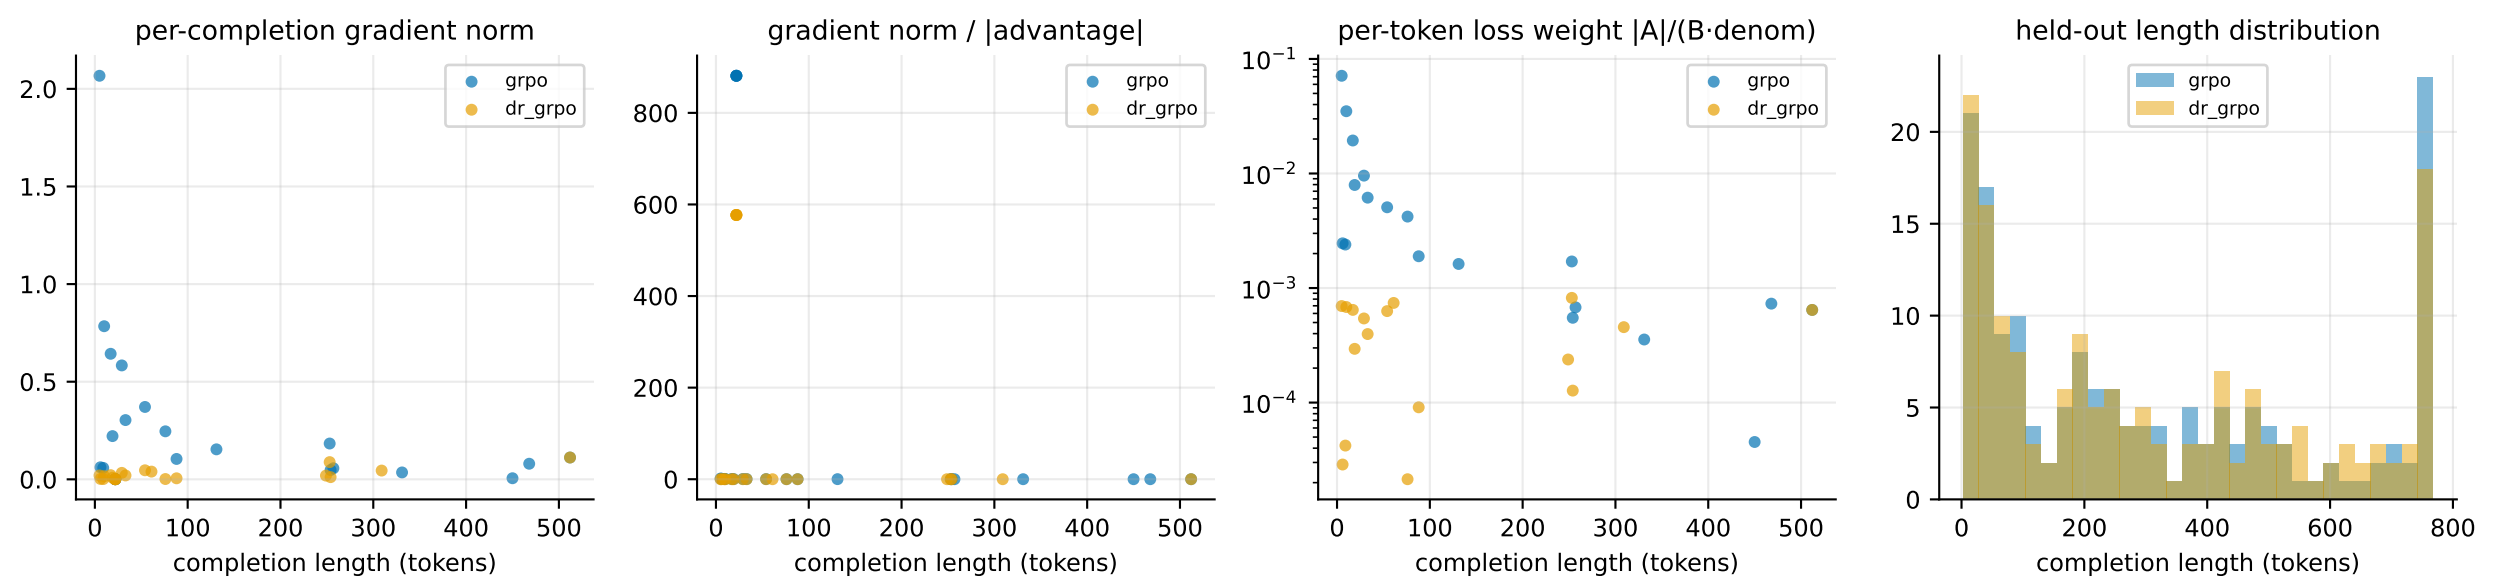
per-completion gradient norm
grpo
dr_grpo
completion length (tokens)
gradient norm / |advantage|
grpo
dr_grpo
completion length (tokens)
per-token loss weight |A|/(B·denom)
grpo
dr_grpo
completion length (tokens)
held-out length distribution
grpo
dr_grpo
completion length (tokens)